

# Plasma: *Parallel L*anguage for System *M*odeling and *A*nalysis

Brian Kahne, Peter Wilson  
Freescale Semiconductor  
Aseem Gupta, Nikil Dutt  
University of California, Irvine

Freescale Semiconductor Confidential Proprietary. Freescale™ and the Freescale logo are trademarks of Freescale Semiconductor, Inc.  
All other product or service names are the property of their respective owners. © Freescale Semiconductor, Inc. 2004



# The Problem

**How do we quickly and efficiently model a complex SoC so that we can perform a trade-off analysis early in the design cycle?**

- The language must be easy to use, with a minimum of extraneous syntax.
- Must be able to easily express concurrency and help the user avoid common pitfalls.
- It must be possible to map different parts of the program to different aspects of the system, e.g. *“this thread runs on this processor”*.

**Future chips are likely to have many programmable cores:**

- How will users develop software for them?
- How will users map their complex application to an existing product?

# Proposal

## These two problems are closely related:

- Both require the development of explicitly parallel programs.
- In both cases, users want a clean way to express this parallelism.

## The goal of the Plasma project is to investigate whether a language with the *appropriate* constructs might be used to ease the task of system modeling and parallel application development:

- Increase productivity through clearer representation.
- Increase productivity and quality through increased compile-time checking of the more difficult-to-get-right aspects of systems models (the concurrency)

## If successful, the language will have wide applicability:

- Useful throughout the life cycle of a design, from initial product definition to software development for the design.

# Existing Work

***SystemC* attempts to handle modeling but does not truly understand parallelism: No help with software development. It is also very hardware-centric.**

***SpecC* is a language with true parallelism, but is still very hardware centric.**

***HandelC* makes it easy to create hardware using a C-like language, but restricts the language, thus making it not relevant for general software development.**

***OpenMP* adds parallel extensions to C++ but is entirely software oriented: It only handles shared memory systems.**

# The Plasma Language

**Plasma is a set of extensions to C++.**

**The basis for Plasma is Occam:**

- **Based on CSP (Communicating Sequential Processes).**
- **Threads are explicitly created by the user and communicate via typed channels.**
- **Plasma adds support for shared variables protected by mutex code and garbage collection (Boehm collector).**


**Plasma adds a discrete-event simulation API useful for modeling hardware.**

**The language can be used in two ways:**

- **For applications development, it provides a convenient means for exploiting the parallelism in a multi-core, multi-threaded system.**
- **For modeling, the parallelism features are used to model the parallelism inherent in the design being modeled.**

# Thread Creation

## To launch a thread:



## Or:

```
// This launches foo as a thread.  
Result<int> res = spawn(foo(1,2,3));  
// This causes us to wait until foo finishes.  
cout << "The result is: " << res.value() << endl;
```

**A replicating form of *par* also exists: The *pfor* block.**

# Channels

**The primary means for communicating between threads is through the use of typed channels.**

**Channels are simply C++ templates, so new types of channels can be added without changing the language. Plasma currently defines the following channels:**

- **Single item channel:** A read blocks if the channel is empty and a write blocks if the channel has data.
- **Queued channel:** The queue may be fixed in size or infinite.
- **Clocked channel:** Writes may occur at any time, but reads are limited to clock boundaries.
- **Time-out channel.**
- **Spawn-result adapter channel.**

**By default multiple threads may write to a channel, but only one thread may read from the channel.**

**This behaviour can be changed by the user.**

# Channels

**Channels are simply C++ templates. To declare a channel for transmitting integers:**

```
Channel<int> chan;
```

**To send data on a channel:**

```
chan.write(10);
```

**To read data from a single channel:**

```
int x = chan.get();
```


**Channel semantics:**

- **Writing to a full channel causes the thread performing the write to block.**
- **Reading from an empty channel causes the thread performing the read to block.**
- **The result is implicit flow-control: Threads do not need to check the state of the channels, they simply perform their operations and will block if needed.**



# Channels

```
int pMain(int argc,  
          const char *argv[])  
{  
    Channel<int> c;  
    par {  
        t1(c);  
        t2(c);  
    }  
    return 0;  
}  
void t2(Channel<int> &c) {  
    c.write(10);  
    c.write(20);  
}  
void t1(Channel<int> &c) {  
    cout << c.read() << '\n';  
    cout << c.read() << '\n';  
}
```



# Channels and the *Alt* Block

**Plasma uses the *alt* block to allow a thread to read from more one than channel. Upon entry to an *alt* block:**

- If only one channel is ready, that channel does a read.
- If more than one channel is ready, a non-deterministic selection is made.
- If no channels are ready, the thread blocks until one of the channels has data.

***Prialt* and *priafor* are similar to *alt*, except that they ensure the ordering specified by the user.**

**All of these blocks may be nested.**

# Channels and the *Alt* Block

## To read from multiple channels:

```
alt {  
    c1.port(int x) {  
        cout << "Read " << x << " from c1.\n";  
    }  
    c2.port(int x) {  
        cout << "Read " << x << " from c2.\n";  
    }  
}
```

Data mapped to variable on a read.

Channels

Code executed when the read occurs.

## A looping form for *alt* also exists:

```
afor (int i = 0; i != channels.size(); ++i) {  
    channels[i].port(int x) { ... }  
}
```

# Processors

## A *Processor* object exists to group together threads.

They allow the user to partition a design based upon hardware resources.

## Threads can be directed to run on specific processors:

- With **par**:

```
par {  
    on (processor1) { ... }  
    on (processor2) { ... }  
}
```

- With **spawn**:

```
processor1.spawn(foo(1,2,3));
```

## Processors can be made to share their ready queue, enabling SMP modeling:

```
Processor p2 = make_sharedproc(processor1);
```

# Priorities

## Threads have priorities.

- The number of priorities can be set by the user.
- Higher priority threads execute before lower priority threads.
- Lowest-priority threads are time-sliced.

## Changing priorities:

- A thread inherits its priority from its parent.
- The priority can be changed by an API call:  
`pSetPriority(0);`
- The priority can be set at thread creation time:

```
par {  
    on (pCurProc(), 4) { ... }  
    on (p2, 5) { ... }  
}
```

or

```
p1.spawn(foo(1, 2, 3), 10);
```

Specifies the  
priority.

# Shared Data Structures

**To create a shared data structure:**

```
pMutex class Foo { ... };
```

**All public member functions, excluding constructors and destructors, will be wrapped with mutual-exclusion code.**

**Channels are implemented using *pMutex*.**

# Simulation Time

**Plasma includes a discrete-event time model, implemented as an API.**

Its presence is optional; support might not exist for a version optimized for software development.

**Two important functions:**

- **pDelay(p<sub>time</sub>\_t):** Thread is idle for specified amount of time.
- **pBusy(p<sub>time</sub>\_t):** Processor is busy computing for the specified amount of time.

**Time only progresses by making calls to these functions.**

**The idea is that software written in Plasma can be annotated in order to understand performance.**

# Power Modeling

**Power is modeled in a similar manner to time: Energy is stored on a per-processor basis.**

- **pEnergy(energy\_t):** Adds energy to the processor.
- **pGetEnergy():** Returns the processor's energy and clear the value.

**To calculate power:**

- **Sample a processor's energy on a periodic basis.**
- **Divide the energy by the sample period.**



# Roadmap

**Currently, we are conducting an experiment to see if this concept makes sense.**

- **We have a front-end which parses Plasma and generates C++.**
- **It targets a simple user-mode threading library implemented with Quickthreads (all threads run in a single UNIX process).**

**We are in the process of releasing Plasma as open-source.**

# Examples

## Several examples have been developed in order to demonstrate Plasma's use:

- **An engine controller to demonstrate Plasma's use for embedded applications.**
- **A transaction processing example: Clients send database requests to a mainframe which might have to send requests to a disk array.**
- **A simple RISC pipeline. This demonstrates the use of clocked channels.**
- **A parallel C compiler: Code generation for functions happens in parallel.**
- **The 2-D Discrete Wavelet Transform (DWT) block used in the JPEG2000 and MPEG4 compression standards.**

# Results

## SystemC vs. Plasma:

- **The DWT block was originally modeled in SystemC, then converted to Plasma.**
- **The biggest advantages that we saw in using Plasma were:**
  - The ease with which threads can be created:
    - > No need to declare classes, mark methods as `SC_THREAD`, etc.
    - > Instead, threads can be launched as functions.
    - > SystemC 2.1 adds some of this capability with `sc_spawn`, but it still requires extra syntax for declaring a class and passing parameters.
  - Communication:
    - > The Plasma code was shorter due to the simpler flow-control mechanisms provided by Plasma's channels.
    - > Plasma's simple channels, since they are templated, were used for a wide variety of communication tasks.
- **Of course, this is all qualitative: Plasma is in too early of a stage of development to do a fair comparison on performance.**

# Conclusion

**Models of complex systems and multi-threaded software both share a common need: An easy way to express explicit parallelism.**

**Our proposal is the *Plasma* language:**

- **Superset of C++.**
- **Parallelism based upon CSP.**
- **Inherent parallelism provides for the potential to optimize context switches, catch common parallel-programming mistakes, etc.**

**The language may be used for both application development and for modeling:**

- **The language provides a minimal set of parallel concepts.**
- **A library provides various types of channels for communication and a discrete-event simulation API.**
- **This allows a problem to first be modeled as a multi-threaded software application, then be decomposed into a hardware/software systems model.**